

Chelimela Raghavendra Goud

Hyderabad, India — +91-7780116630 — raghava1402.ch@gmail.com

github.com/raghavendra1422 — linkedin.com/in/raghavendra1422

PROFESSIONAL SUMMARY

Full stack Developer-GenAI, with hands-on experience building production RAG pipelines and agentic workflows using LangChain, LangGraph, CrewAI, FastAPI, and PostgreSQL. Delivered 80% compliance decision accuracy on a self-built Agentic RAG system and shipped AI-powered platforms integrating OpenAI, Gemini, and Grok APIs. Actively transitioning into a dedicated AI Engineer role, with strong foundations in semantic search, prompt engineering, hallucination detection, multi-agent orchestration, and LLM evaluation.

CORE TECHNICAL SKILLS

Programming & Frontend: Python, JavaScript, React.js, TSX, HTML, CSS

Backend & Databases: FastAPI, REST APIs, Uvicorn, Swagger, PostgreSQL, ChromaDB, Vector Databases

Generative AI: LLMs, RAG, Prompt Engineering, Semantic Search, Embeddings, AI Agents, Re-ranking, Tool Calling

Frameworks & Tools: LangChain, CrewAI, n8n, Hugging Face, RAGAS, DeepEval

DevOps: Git, GitHub, Docker, CI/CD Concepts

Integrations: OpenAI API, Gemini API, xAI Grok API, YouTube Data API, BeautifulSoup, Web Scraping, Data Extraction Pipelines

PROFESSIONAL EXPERIENCE

Full stack Developer-GenAI

Aug 2025 – Present

GenAI Lakes Pvt Ltd, Hyderabad

- Architected and shipped 3 production AI workflow systems — keyword intelligence, competitor monitoring, and digital presence analysis — reducing manual research time by ~60% for end users.
- Built and deployed 10+ FastAPI REST endpoints backed by PostgreSQL, serving real-time dashboard analytics and structured workflow orchestration with <200ms p95 latency under concurrent load.
- Integrated 3 LLM providers (OpenAI, Gemini, Grok) into a unified semantic retrieval and recommendation pipeline, enabling multi-model fallback and cutting per-query cost by ~30% through provider routing.
- Designed 5+ agentic workflows using CrewAI and n8n, automating multi-step competitor analysis and lead enrichment tasks that previously required 4-6 hours of manual effort per cycle.

AI Prompt Engineer & Evaluator

Nov 2024 – Mar 2025

Outlier (Remote)

- Evaluated 500+ LLM responses across reasoning, hallucination detection, and contextual accuracy tasks; developed prompt optimization strategies that improved output quality scores by ~25% on assigned benchmarks.

PROJECTS

Compliance-Guard AI – RBI Loan Compliance Checker

RAG, LangChain, FastAPI, ChromaDB, Groq LLM, React, DeepEval, Python

- Built a production Agentic RAG system achieving 80% end-to-end compliance accuracy on RBI regulation checks; implemented hallucination detection with confidence-based retry loops, reducing false-positive compliance approvals by ~35%.
- Engineered automated PDF version control replacing stale ChromaDB chunks on new circular uploads — maintaining 100% up-to-date regulatory knowledge — with async FastAPI endpoints and a React audit dashboard for real-time compliance reporting.

Deep Research AI Agent – Autonomous Multi-Step Research Pipeline

LangGraph, Groq LLM, Serper API, HuggingFace Embeddings, ChromaDB, Python

- Designed a 5-node LangGraph state machine (Planner → Researcher → Extractor → Writer → Memory) that autonomously breaks a query into sub-topics, performs real-time web research via Serper API, and synthesizes a structured markdown report.
- Implemented persistent semantic memory using ChromaDB with local HuggingFace embeddings, enabling instant follow-up Q&A on generated reports without repeated LLM calls or API cost.

EDUCATION

Bachelor of Technology (B.Tech) – Electronics and Communication Engineering

CMR Engineering College, Hyderabad

CGPA: 8.02