

Chelimela Raghavendra Goud

Hyderabad, India — +91-7780116630 — raghava1402.ch@gmail.com

github.com/raghavendra1422 — linkedin.com/in/raghavendra1422

PROFESSIONAL SUMMARY

AI Engineer with hands-on experience building production RAG pipelines, agentic workflows, and LLM-powered automation systems using LangChain, CrewAI, FastAPI, and PostgreSQL. Delivered 80% compliance decision accuracy on a real-world Agentic RAG system and shipped 3 production AI platforms integrating OpenAI, Gemini, and Grok APIs. Experienced in semantic search, prompt engineering, hallucination detection, multi-agent orchestration, and LLM evaluation — focused on building AI systems that work reliably at scale.

CORE TECHNICAL SKILLS

Programming & Frontend: Python, JavaScript, React.js, TSX, HTML, CSS

Backend & Databases: FastAPI, REST APIs, Uvicorn, Swagger, PostgreSQL, ChromaDB, Vector Databases

Generative AI: LLMs, RAG, Prompt Engineering, Semantic Search, Embeddings, AI Agents, Re-ranking, Tool Calling

Frameworks & Tools: LangChain, CrewAI, n8n, Hugging Face, RAGAS, DeepEval

DevOps: Git, GitHub, Docker, CI/CD Concepts

Integrations: OpenAI API, Gemini API, xAI Grok API, YouTube Data API, BeautifulSoup, Web Scraping, Data Extraction Pipelines

PROFESSIONAL EXPERIENCE

AI Engineer

Aug 2025 – Present

GenAI Lakes Pvt Ltd, Hyderabad

- Architected and shipped 3 production AI workflow systems — keyword intelligence, competitor monitoring, and digital presence analysis — reducing manual research time by ~60% for end users.
- Built and deployed 10+ FastAPI REST endpoints backed by PostgreSQL, serving real-time dashboard analytics and structured workflow orchestration with <200ms p95 latency under concurrent load.
- Integrated 3 LLM providers (OpenAI, Gemini, Grok) into a unified semantic retrieval and recommendation pipeline, enabling multi-model fallback and cutting per-query cost by ~30% through provider routing.
- Designed 5+ agentic workflows using CrewAI and n8n, automating multi-step competitor analysis and lead enrichment tasks that previously required 4–6 hours of manual effort per cycle.

AI Prompt Engineer & Evaluator

Nov 2024 – Mar 2025

Outlier (Remote)

- Evaluated 500+ LLM responses across reasoning, hallucination detection, and contextual accuracy tasks; developed prompt optimization strategies that improved output quality scores by ~25% on assigned benchmarks.

PROJECTS

FastLeads99 – AI-Powered Competitive Intelligence Platform

React.js, TypeScript, FastAPI, PostgreSQL, OpenAI, Gemini, Grok, YouTube API

- Built a full-stack AI intelligence platform processing 1,000+ keywords/run across competitor analysis, SEO trend monitoring, and lead generation — reducing manual research cycles from days to under 10 minutes.
- Developed FastAPI backend services, PostgreSQL data models, and AI-driven keyword intelligence systems for workflow automation and dashboard analytics.
- Integrated YouTube APIs, Google APIs, and BeautifulSoup-based scraping pipelines for SEO insights, competitor data extraction, and trending keyword analysis.

Compliance-Guard AI – RBI Loan Compliance Checker

LangChain, FastAPI, ChromaDB, Groq LLM, React, DeepEval, Python

- Built a production Agentic RAG system achieving 80% end-to-end compliance accuracy on RBI regulation checks; implemented hallucination detection with confidence-based retry loops, reducing false-positive compliance approvals by ~35%.
- Engineered automated PDF version control replacing stale ChromaDB chunks on new circular uploads — maintaining 100% up-to-date regulatory knowledge — with async FastAPI endpoints and a React audit dashboard for real-time compliance reporting.

EDUCATION

Bachelor of Technology (B.Tech) – Electronics and Communication Engineering

CMR Engineering College, Hyderabad

CGPA: 8.01